

# CS 305: Computer Networks

## Fall 2022

### **Lecture 4: Application Layer**

**Ming Tang**

Department of Computer Science and Engineering  
Southern University of Science and Technology (SUSTech)

# HTTP Outline

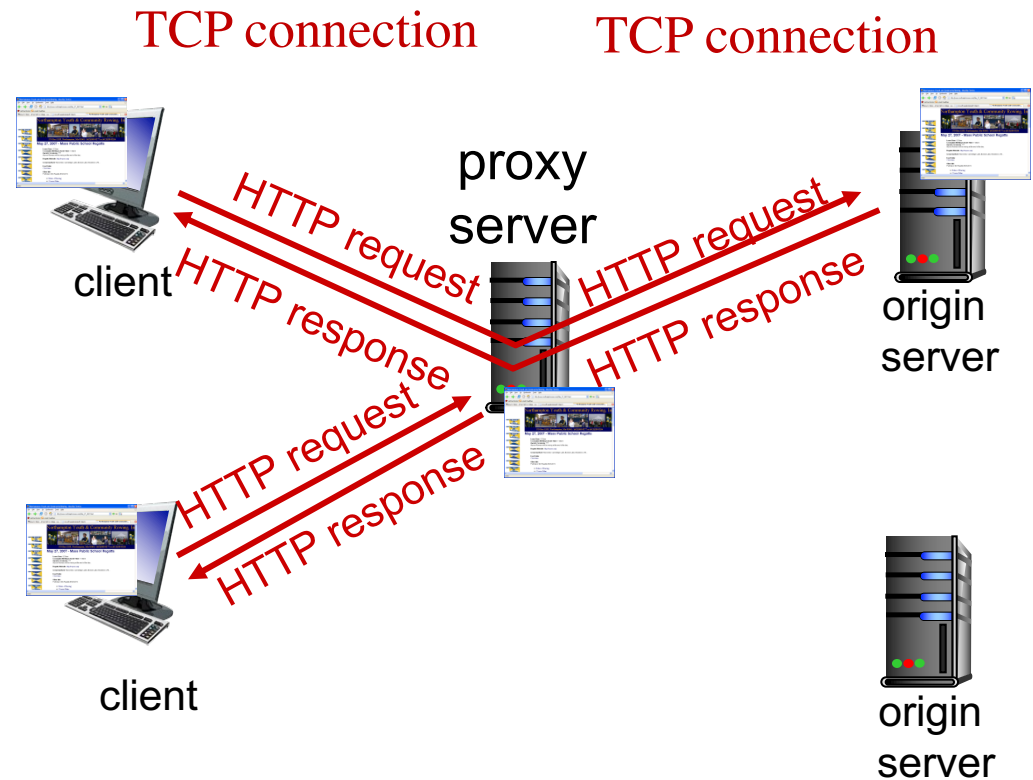
- HTTP Overview
  - HTTP runs over TCP
  - HTTP is stateless
  - Persistent and non-persistent connection
- Request and response messages
- Cookies
- Web caching

# Web caches: proxy (代理) server

*goal:* satisfy client request without involving origin server

Browser sends all HTTP requests to cache

- object in cache: cache returns object
- else cache requests object from origin server, then returns object to client



# More about Web caching

- Cache (Proxy server) acts as both client and server
  - server for original requesting client
  - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)

## *Why Web caching?*

- reduce response time for client request (bottleneck bandwidth)
- reduce traffic on an institution's **access link**
- Internet dense with caches: enables “poor” **content providers** to effectively deliver content (so too does P2P file sharing)


# Caching example:

## Assumptions:

- avg object size: 1M bits
- avg request rate from browsers to origin servers: 15 requests/sec
- avg data rate to all browsers: 15 Mbps
- RTT from router A to any origin server: 2 sec  
→ “Internet delay”
- access link rate: 15.4 Mbps

## Consequences:

- LAN utilization:  $15\text{Mbps}/100\text{Mbps}=0.15$
- access link utilization =  $15/15.4=0.974$
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + milliseconds




# Caching example: fatter access link

## *assumptions:*

- avg object size: 1M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 15 Mbps
- RTT from router A to any origin server: 2 sec
- access link rate: 15.4 Mbps

## *consequences:*

- LAN utilization: 0.15
- access link utilization = ~~0.974~~ → 0.1
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + ~~minutes~~ → milliseconds



*Cost:* increased access link speed (not cheap!)

# Caching example: install local cache

## *assumptions:*


- avg object size: 1M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 15 Mbps
- RTT from router A to any origin server: 2 sec
- access link rate: 15.4 Mbps

## *consequences:*

- LAN utilization: 0.15
- access link utilization = ?
- total delay = ?

*How to compute link utilization, delay?*

**Cost:** web cache (cheap!)




Hit rates: the fraction of requests that are satisfied by a cache. Typically, 0.2—0.7.

# Caching example: install local cache

*Calculating access link utilization, delay with cache:*


- suppose cache hit rate is 0.4
  - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
  - 60% of requests use access link
- data rate to browsers over access link  
 $= 0.6 * 15 \text{ Mbps} = 9 \text{ Mbps}$ 
  - utilization  $= 9 / 15.4 = 0.58$
- Average delay
  - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
  - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
  - less than with 150 Mbps link (and cheaper too!)



Typically, a traffic intensity less than 0.8 corresponds to a small delay, say, tens of milliseconds



# Conditional GET



The copy of an object residing in the cache may be **out-of-date**:

## Conditional GET

- GET method
- If-Modified-Since

```
GET /fruit/kiwi.gif HTTP/1.1
Host: www.exotiquecuisine.com
If-modified-since: Wed, 9 Sep 2015 09:23:24
```

**Goal:** allows a cache to verify that its objects are **up to date**

- don't send object if cache has up-to-date cached version
- no object transmission delay
- lower link utilization

# Conditional GET

When a browser requests an object via proxy cache:

Proxy  
cache



server

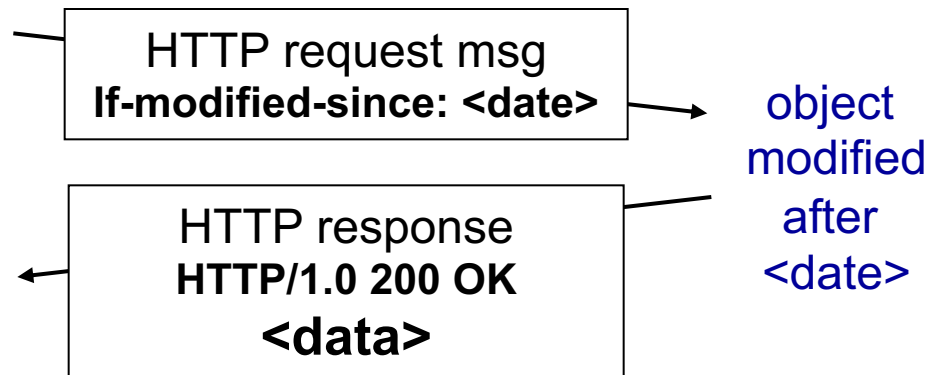
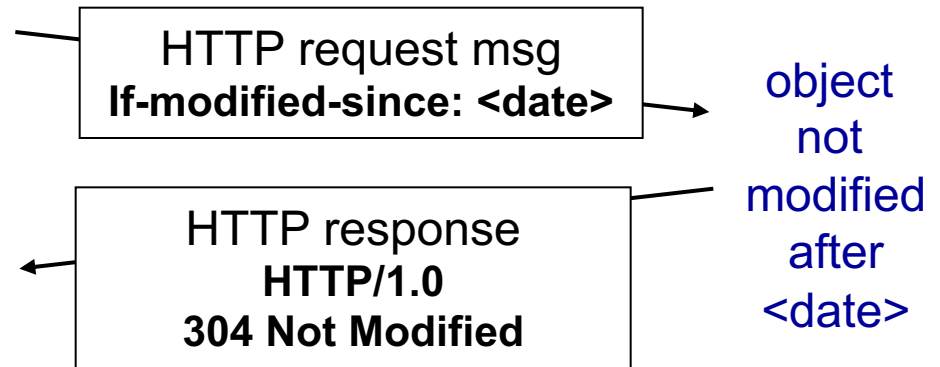


- *Proxy cache*: specify date of cached copy in HTTP request  
**If-modified-since: <date>**
- *Server*: response contains no object if cached copy is up-to-date:

## HTTP/1.0 304 Not Modified

```
HTTP/1.1 304 Not Modified
Date: Sat, 10 Oct 2015 15:39:29
Server: Apache/1.3.0 (Unix)
```

*(empty entity body)*



# HTTP Summary

- HTTP Overview
  - HTTP runs over TCP
  - HTTP is stateless
  - Persistent and non-persistent connection
- Request and response messages
- Cookies
- Web caching

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

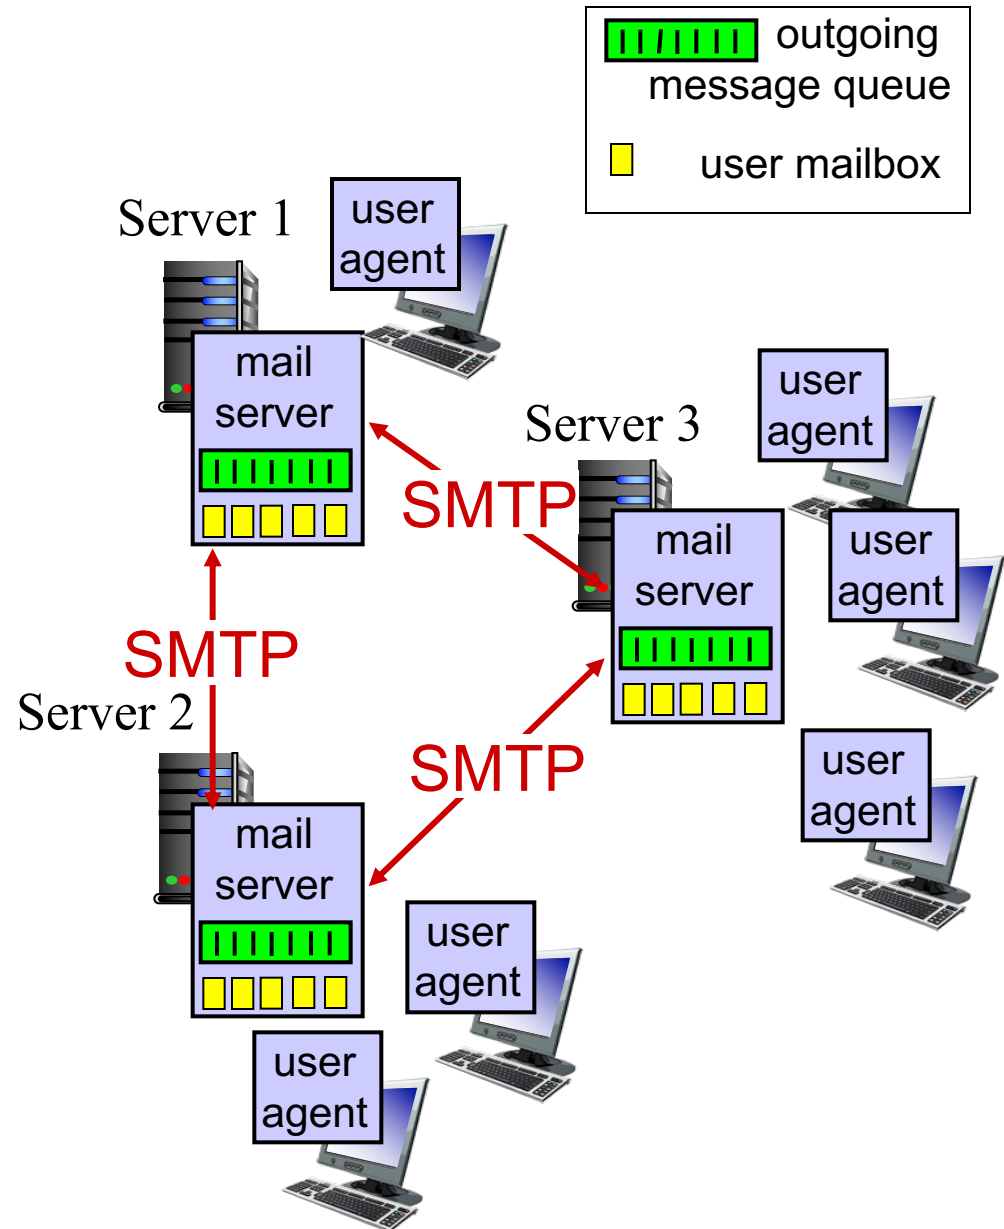
# Electronic Mail Overview

- Overview
  - Main components
  - Alice sends an email to Bob
- SMTP
- Mail Message Format
- Mail Access Protocol
  - POP3
  - IMAP
  - HTTP: Web-based Email

# Electronic mail

## Three major components:


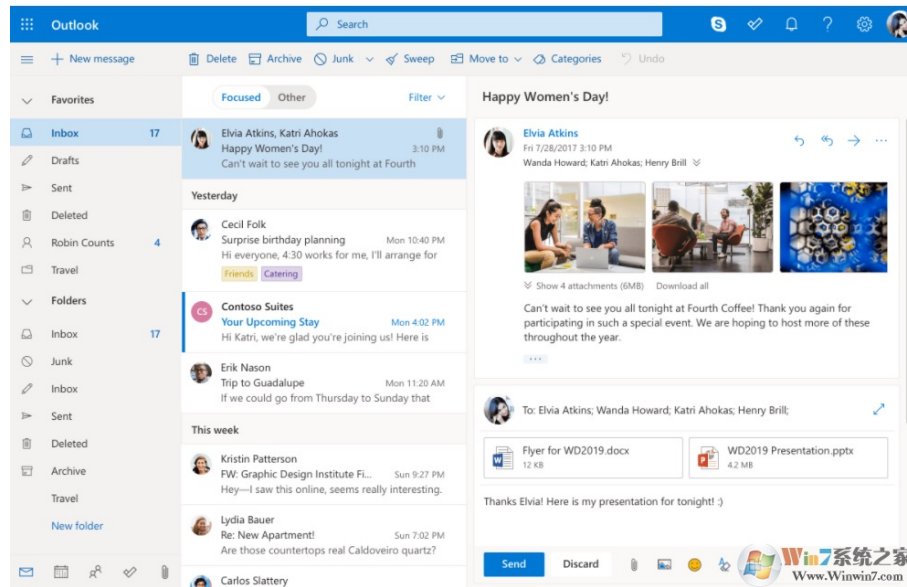
- user agents
- mail servers
- simple mail transfer protocol (SMTP): use TCP



# Electronic mail: User Agent

## User Agent


- a.k.a. “mail reader”
- Allow users to read, reply to, forward, save and compose messages
- e.g., Outlook, iPhone mail client



# Electronic mail: mail servers


## Mail servers:

- Always-on hosts
  - *User mailbox* contains outgoing, incoming messages
  - *Message queue* of outgoing (to be sent) mail messages
  - *Simple Mail Transfer Protocol (SMTP)* between mail servers to send email messages
    - client: sending mail server
    - “server”: receiving mail server
- Both client and server sides of SMTP run on mail server.

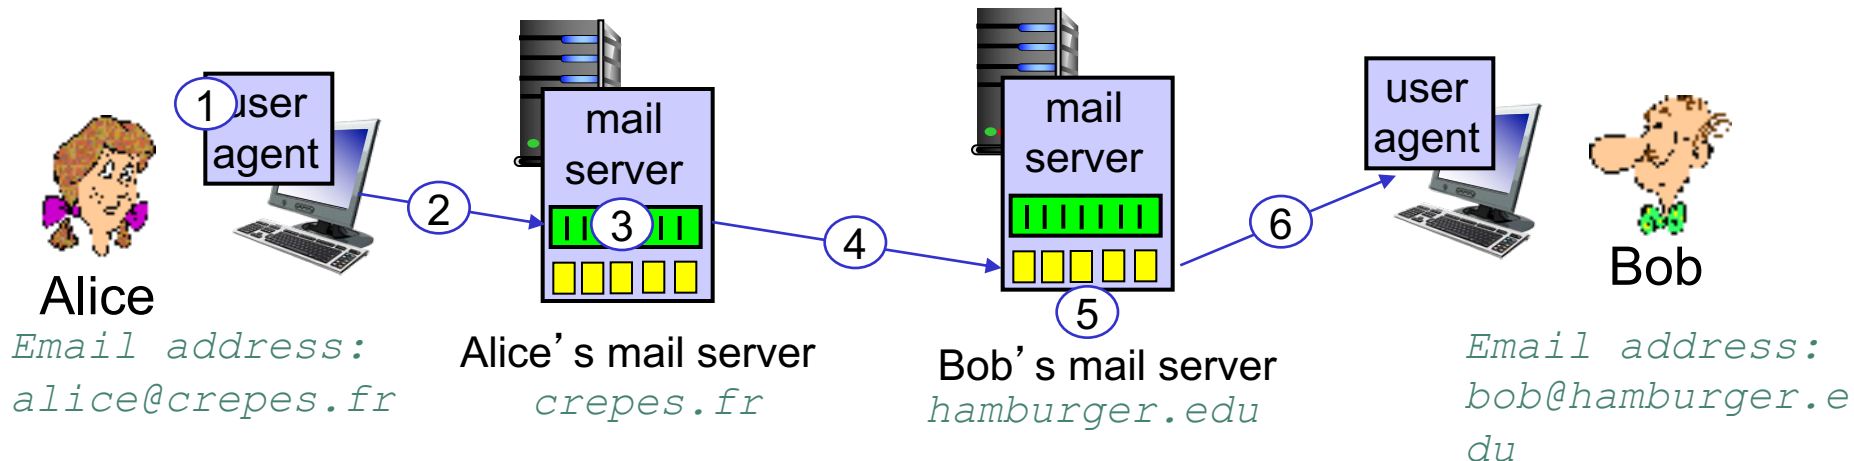




# Scenario: Alice sends message to Bob



- 1) Alice uses user agent to compose message "to" bob@hamburger.edu
- 2) Alice's user agent sends message to her mail server; message placed in message queue
- 3) **client side** of SMTP opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's message over the TCP connection
- 5) Bob's mail server places the message in Bob's mailbox
- 6) Bob invokes his user agent to read message

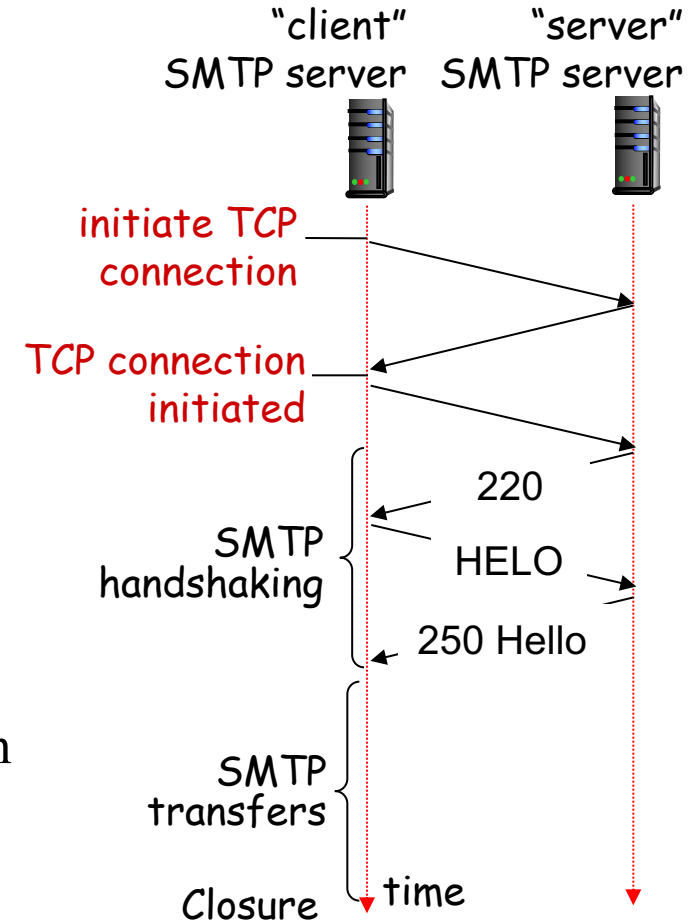


# Electronic Mail Overview

- Overview
  - Main components
  - Alice sends an email to Bob
- SMTP
- Mail Message Format
- Mail Access Protocol
  - POP3
  - IMAP
  - HTTP: Web-based Email

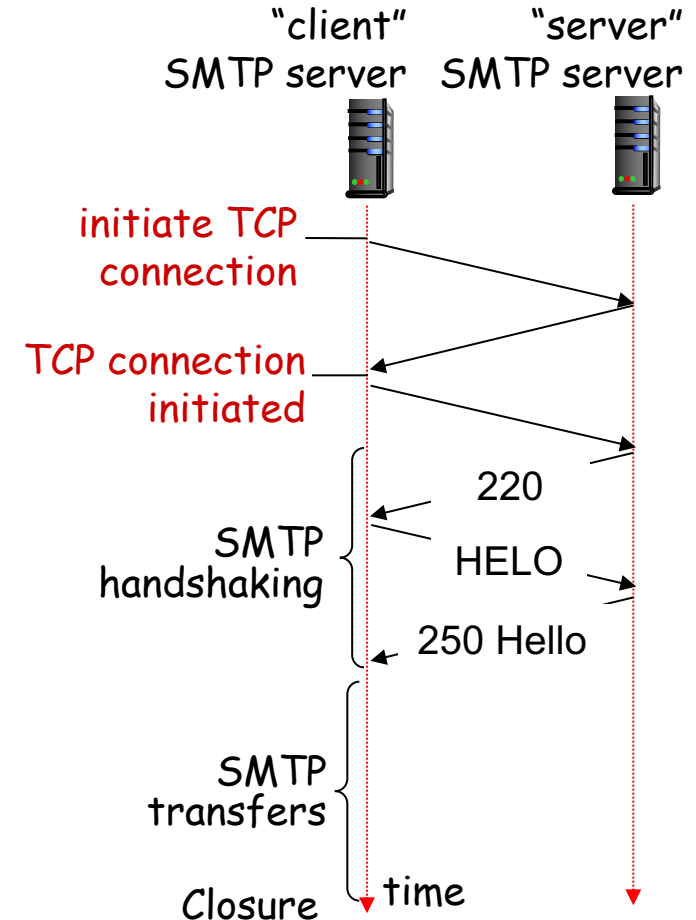
# Electronic Mail: SMTP [RFC 2821]

- Uses TCP to reliably transfer email message from client to server, port 25
  - If fail, new attempt after a while (e.g., 30 minutes)
- Direct transfer: sending server to receiving server
  - Direct connection, no intermediate mail server
- Three phases of transfer
  - handshaking (greeting): indicate email address
  - transfer of messages: persistent connection
  - closure



# Electronic Mail: SMTP [RFC 2821]

- Two types of messages (like HTTP)
  - **commands:** text
  - **response:** status code and phrase
- Entire messages (header & body) must be in ASCII
  - Binary multimedia data → ASCII
  - For HTTP, headers are encoded with ASCII



# Sample SMTP interaction

The following are exactly the lines the client (C: crepes.fr) and server (S: hamburger.edu) send after they establishing TCP connections.

		<b>commands</b> response (status code + phrase)
SMTP handshaking	S:	220 hamburger.edu
	C:	HELO crepes.fr
	S:	250 Hello crepes.fr, pleased to meet you
SMTP transfers	C:	MAIL FROM: <alice@crepes.fr>
	S:	250 alice@crepes.fr... Sender ok
	C:	RCPT TO: <bob@hamburger.edu>
	S:	250 bob@hamburger.edu ... Recipient ok
	C:	DATA
	S:	354 Enter mail, end with "." on a line by itself
	C:	Do you like ketchup?
	C:	How about pickles?
	C:	.
	S:	250 Message accepted for delivery
Closure	C:	QUIT
	S:	221 hamburger.edu closing connection

Repeat to send multiple messages

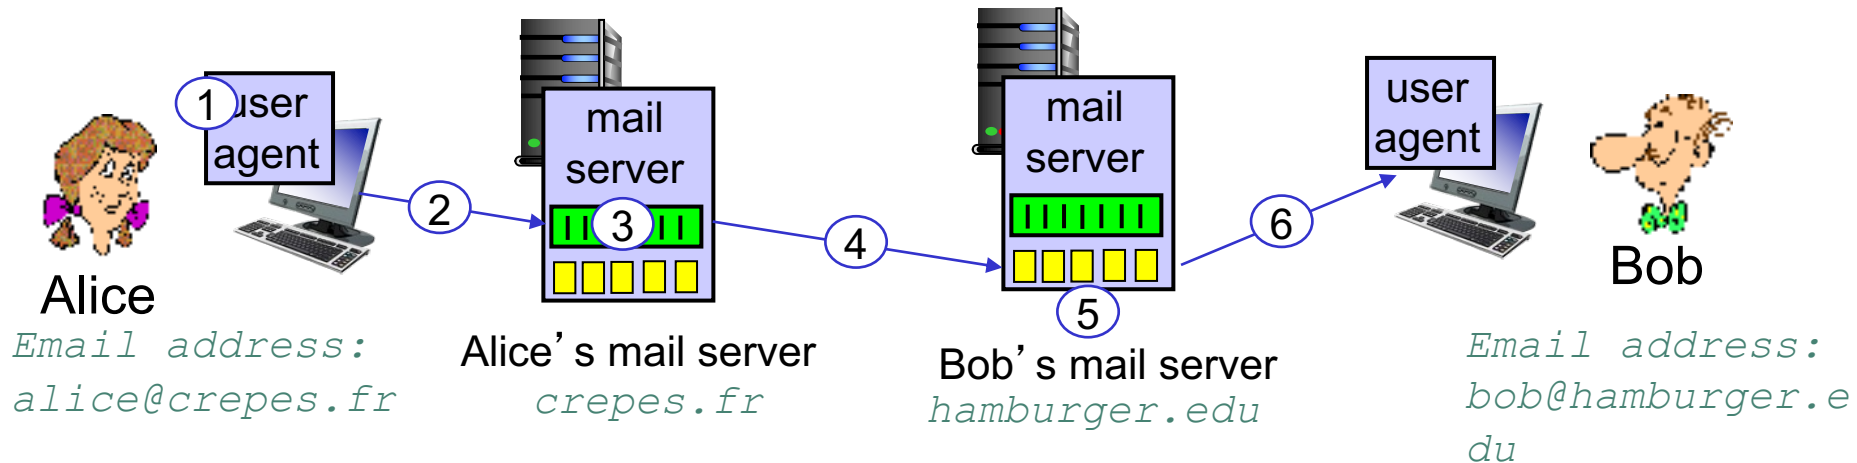
# SMTP: Closing Observations

- SMTP uses persistent connections
- SMTP requires message (header & body) to be in ASCII
- SMTP server uses CRLF.CRLF to determine end of message

## Comparison with HTTP:

- HTTP: pull
- SMTP: push
- HTTP: ASCII in header
- SMTP: ASCII in header and body
- HTTP: each object encapsulated in its own response message
- SMTP: multiple objects sent in one message

# Alternative Choices?




Can we have mail servers directly on user's local PC?

**NO**

Can we let Alice send to Bob's mail server directly?

**NO!**

# Alternative Choices?




Why **not** having mail servers directly on user's local PC?

- Recall that a mail server manages mailboxes and runs the client and server sides of SMTP.
- If Bob's mail server were to reside on his local PC, then Bob's PC would have to remain always on in order to receive new mail.



# Alternative Choices?



Why **not** letting Alice send to Bob's mail server directly?

- Bob's mail sever may fail; need to repeatedly send the message until success.

# Electronic Mail Overview


- Overview
  - Main components
  - Alice sends an email to Bob
- SMTP
- Mail Message Format
- Mail Access Protocol
  - POP3
  - IMAP
  - HTTP: Web-based Email

# Mail message format



Mail message format (RFC 2822) defines *syntax* for e-mail message itself (like HTML)

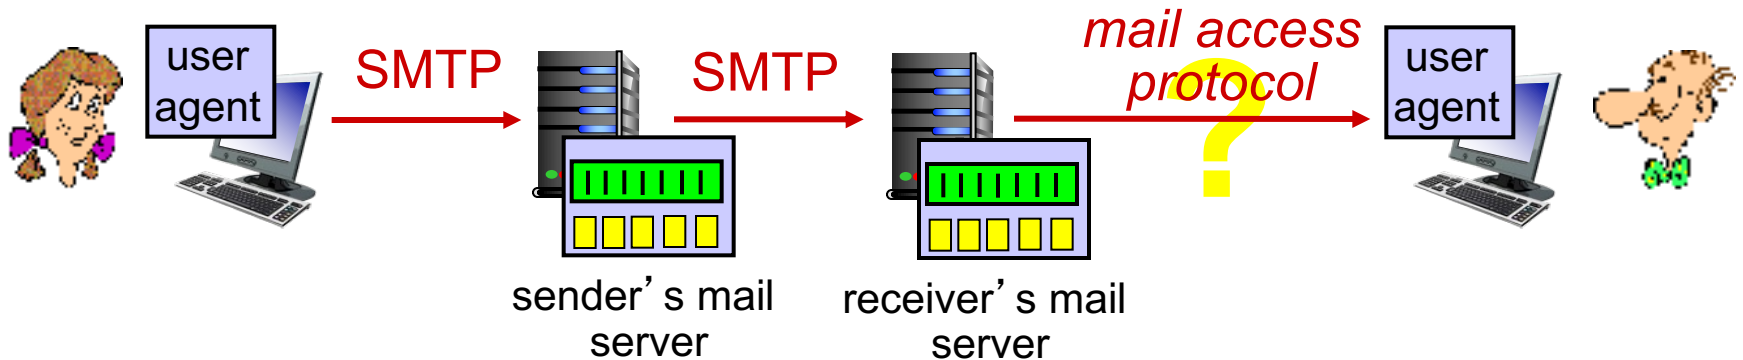
- Header lines, e.g.,
  - To:
  - From:
  - Subject:
  - these lines are part of the message itself, different from SMTP MAIL FROM:, RCPT TO: commands!
- Body: the “message”, ASCII characters only



# Electronic Mail Overview

- Overview
  - Main components
  - Alice sends an email to Bob
- SMTP
- Mail Message Format
- Mail Access Protocol
  - POP3
  - IMAP
  - HTTP: Web-based Email

# Mail access protocols



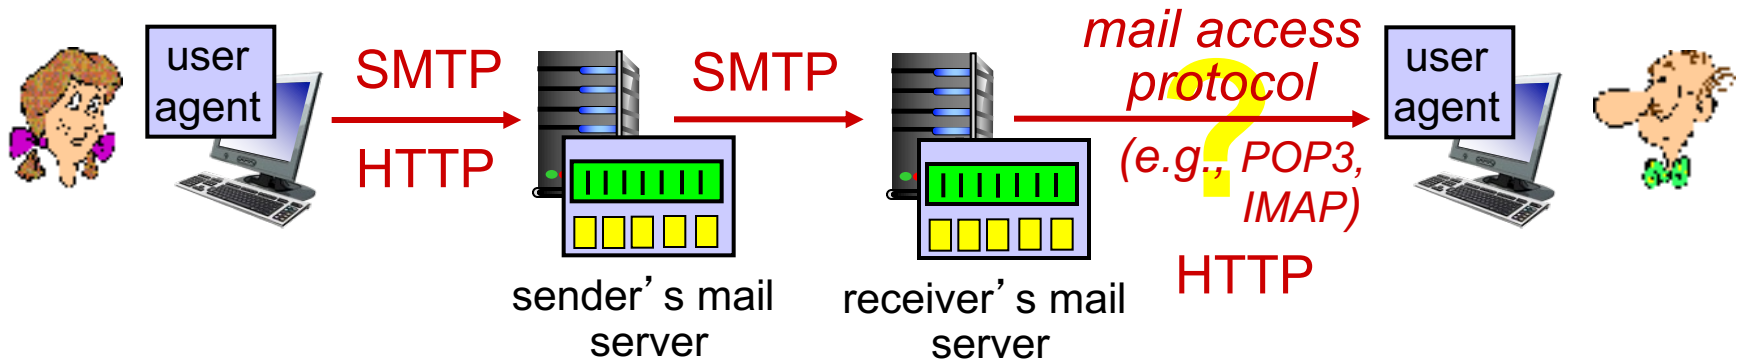
**SMTP:** delivery to receiver's server

**Mail access protocols:** How does Bob obtain his message?

SMTP?

No! Because obtaining message is a **pull operation**.


# Mail access protocols



## **Mail access protocol:** retrieval from server

- **POP3:** Post Office Protocol 3: authorization, download
  - TCP, port 110
- **IMAP:** Internet Mail Access Protocol: more features, including maintain folders, keep user state
- **HTTP:** gmail, Hotmail, Yahoo! Mail, etc.

# POP3 protocol



## Authorization phase

- client commands:
  - user:** declare username
  - pass:** password
- server responses
  - +OK**
  - ERR**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on
```

## Transaction phase

- client:
  - list:** list message numbers
  - retr:** retrieve message by number
  - dele:** delete
  - Quit**

```
C: list
S: 1 498
S: 2 912
S: .
```

**Download-and-delete mode**

```
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
```

```
C: dele 2
C: quit
S: +OK POP3 server signing off
```

## Update phase

- After **Quit**, the mail server deletes the messages marked as deletion

**Download-and-keep mode ?**

# POP3 (more) and IMAP

## More about POP3

- previous example uses POP3 “download and delete” mode
  - Bob cannot re-read e-mail if he changes client
- POP3 “download-and-keep”: reread the message from different machines
- POP3 is stateless across sessions

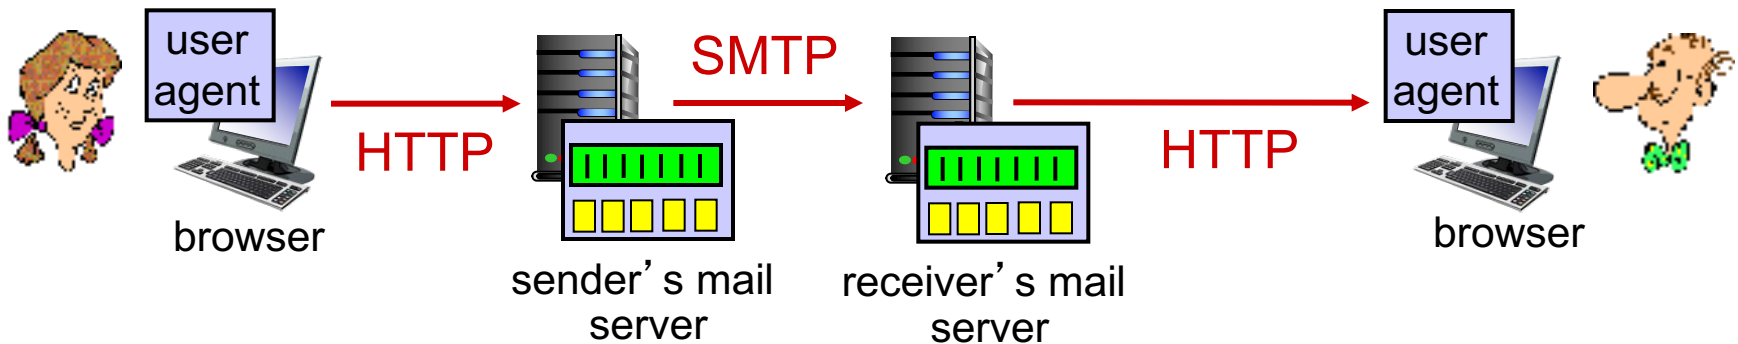
## IMAP

- Maintain a folder hierarchy in one place: at server
- allows user to organize messages in folders
- keeps user state across sessions:
  - names of folders and mappings between message IDs and folder name
- Obtain components of messages





# Web-based Email



Web-based emails are provided by gmail, Hotmail, Yahoo! Mail, etc.

- The user agent is an ordinary web browser

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

# DNS: domain name system

People: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- hostname, e.g.,  
www.yahoo.com -  
used by humans
- IP address (32 bit) -  
used for addressing  
datagrams

Q: how to map between IP  
address and name, and  
vice versa ?

Domain Name System (DNS):

- distributed database implemented  
in hierarchy of many *name  
servers*
- application-layer protocol: hosts  
and name servers communicate  
to *resolve* names (address/name  
translation)

# DNS Overview

- DNS Services
- DNS Structure
  - Hierarchical structure
  - Iterated and recursive query
- DNS protocol
  - DNS Records
  - Query and reply messages
- Inserting records into DNS

# DNS Services

- **hostname to IP address translation**
- **host aliasing**
  - canonical, alias hostnames
  - **www.ibm.com** (alias) is really **servereast.backup2.ibm.com** (canonical)
  - From supplied alias hostname to canonical hostname
- **mail server aliasing**
- **load distribution**
  - replicated Web servers: many IP addresses correspond to one name
  - rotation distributes the traffic (rotate the ordering of IP addresses)



# DNS Services

---


1. An application invokes the client side of DNS
  - specifying the **hostname** that needs to be translated
2. DNS in the user's host takes over, sending a query message into the network.
  - DNS query and reply messages
  - **UDP datagrams** to port 53.
3. After a delay, ranging from milliseconds to seconds, DNS in the user's host receives a DNS reply message that provides the desired mapping.
4. The **mapping (hostname - IP)** is then passed to the invoking application.

## Why UDP?

- fast speed
- smaller data packets

# DNS Services

From the perspective of the invoking application in the user's host, DNS is a **black box** providing a simple, straightforward translation service.



# DNS Overview

- DNS Services
- DNS Structure
  - Hierarchical structure
  - Iterated and recursive query
- DNS protocol
  - DNS Records
  - Query and reply messages
- Inserting records into DNS



# DNS Structure

---

## Centralized DNS:


Clients simply direct all queries to the single DNS server, and the DNS server responds directly to the querying clients.

### *Why not centralize DNS?*

- Single point of failure
- Traffic volume
- Distant centralized database
- Maintenance: huge database, update frequently

A: **doesn't scale!**

# DNS: a distributed, hierarchical database

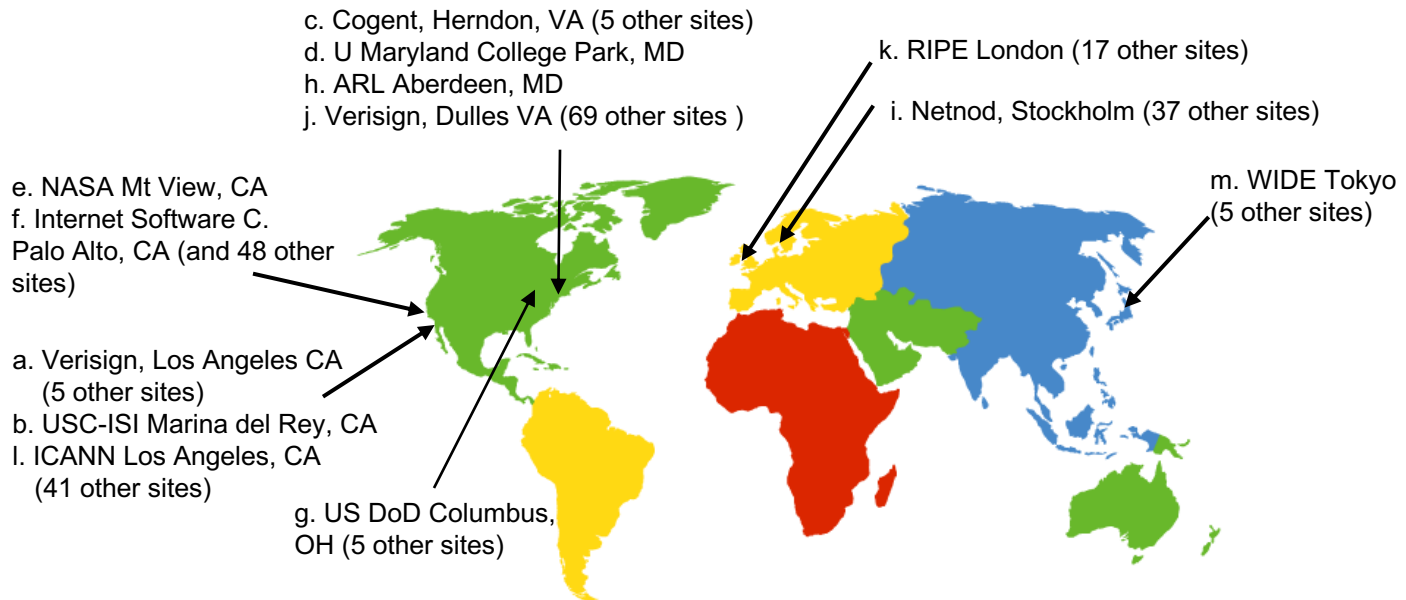


Client wants IP for **www.google.com / scholar.google.com**:

- **Root DNS Servers**: find IP address of the **.com** TLD DNS server
- **Top-Level Domain (TLD) DNS**: client queries **.com** DNS server to get **google.com** authoritative DNS server
- **Authoritative DNS servers**: client queries **google.com** DNS server to get IP address for **www.google.com / scholar.google.com**

# DNS: root servers

- Root name server:
  - Provide the IP addresses of the TLD servers



13 logical root name  
“servers” worldwide

# TLD, authoritative servers

## Top-level domain (TLD) servers:

- Top-level domains: com, org, net, edu, aero, jobs, museums; top-level country domains: uk, fr, ca, jp
- *Network Solutions* maintains servers for .com TLD
- *Educause* for .edu TLD

## Authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

# Local DNS server

- Does not strictly belong to hierarchy
- Each ISP (residential ISP, company, university) has one
  - also called “default name server”

When a host connects to an ISP, the ISP provides the **IP addresses** of one or more of local DNS servers

- A host’s local DNS server may be typically “close to” the host

When host makes DNS query, query is sent to local DNS server

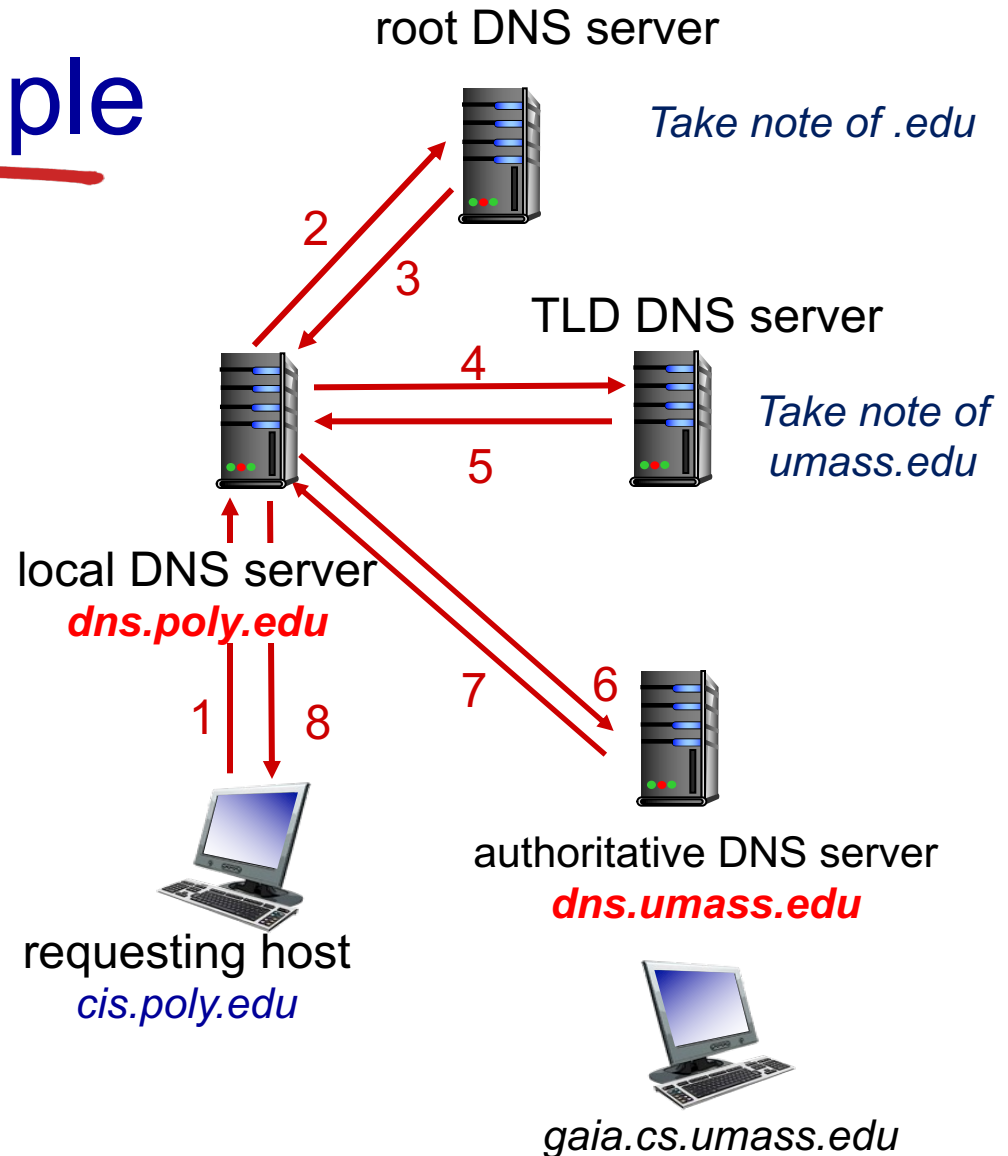
- **acts as proxy**, forwards query into hierarchy
- has **local cache** of recent name-to-address translation pairs (but may be out of date!)

# DNS name resolution example

- host at cis.poly.edu wants IP address for **gaia.cs.umass.edu**

## Iterated query:

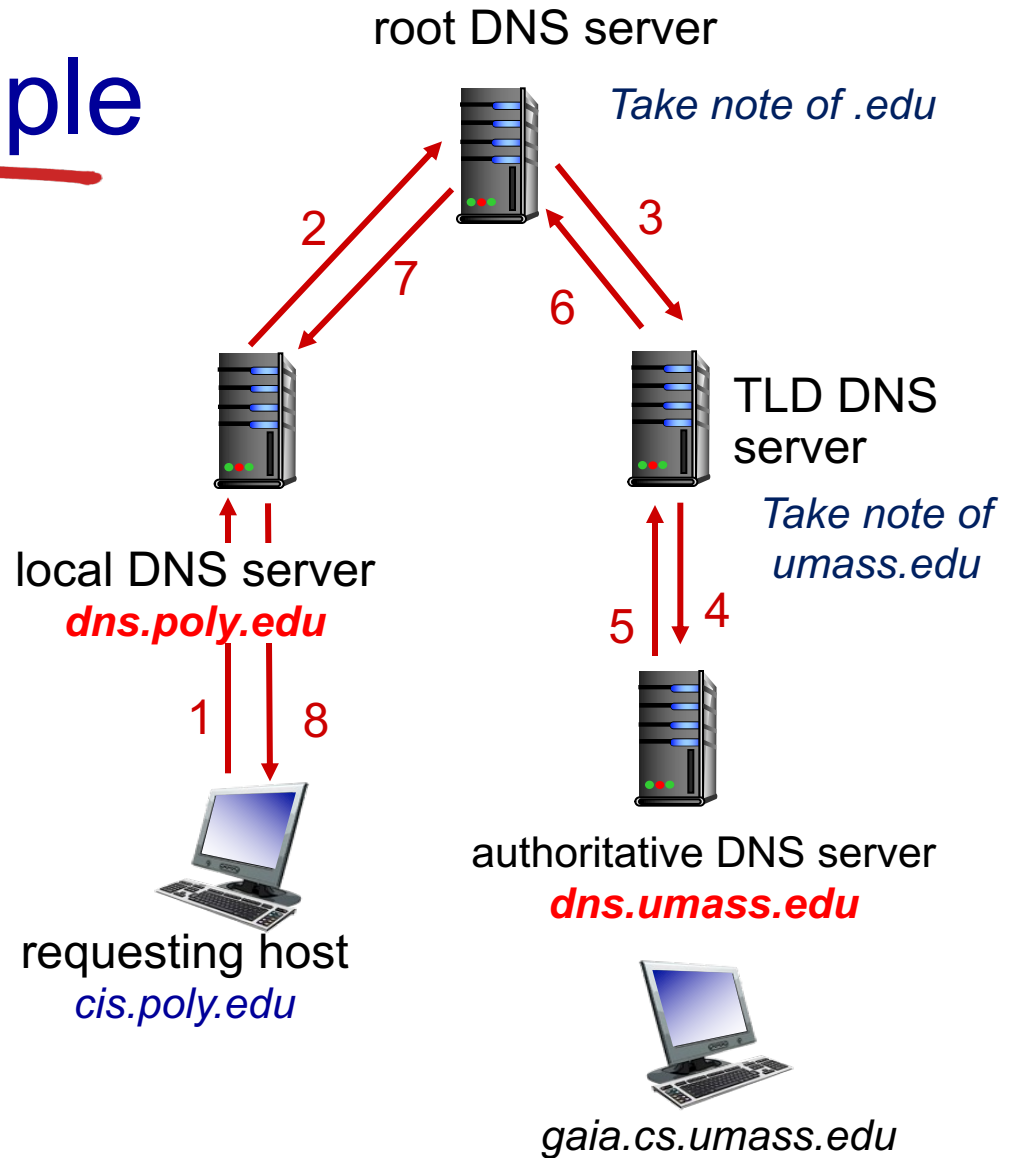
- contacted server replies with the name of another server to contact
- “I don’t know this name, but ask this server”



# DNS name resolution example

## Recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



# DNS: caching, updating records

- Once (any) name server learns mapping, it *caches* mapping
  - TLD servers typically cached in local name servers
  - thus root DNS servers not often visited
- Cached entries may be *out-of-date*
  - cache entries timeout (disappear) after some time (e.g., two days)
- Update/notify mechanisms proposed IETF standard
  - RFC 2136



# DNS Overview

- DNS Services
- DNS Structure
  - Hierarchical structure
  - Iterated and recursive query
- **DNS protocol**
  - DNS Records
  - Query and reply messages
- Inserting records into DNS

# DNS records

**DNS:** distributed database storing resource records (**RR**)

RR format: (name, value, type, ttl)

## type=A

- **name** is hostname
- **value** is IP address

## type=NS

- **name** is domain  
(e.g., foo.com)
- **value** is hostname of  
authoritative server for  
this domain  
(e.g., dns.foo.com)

## type=CNAME

- **name** is alias name for some  
“canonical” (the real) name
- **www.ibm.com** is really  
**servereast.backup2.ibm.com**
- **value** is canonical name

## type=MX

- **value** is canonical name of the  
**mailserver** with **name** (alias name)

# DNS records

If a DNS server is **authoritative** for a particular hostname

- the DNS server will contain a Type A record for the hostname
- (Even if the DNS server is not authoritative, it may contain a Type A record in its cache.)

If a server is **not authoritative for** a hostname

- the server will contain a Type NS record for the domain that includes the hostname
- it will also contain a Type A record that provides the IP address of the DNS server in the **Value** field of the NS record.

Example: an .edu TLD server is not authoritative for gaia.cs.umass.edu


- (umass.edu, dns.umass.edu, NS) .
- (dns.umass.edu, 128.119.40.111, A)

# DNS protocol, messages

Query and reply messages, both with same message format

## message header

- **identification:** 16 bit number for query, reply to query uses same number
- **flags:**
  - query or reply
  - recursion desired
  - recursion available
  - reply is authoritative



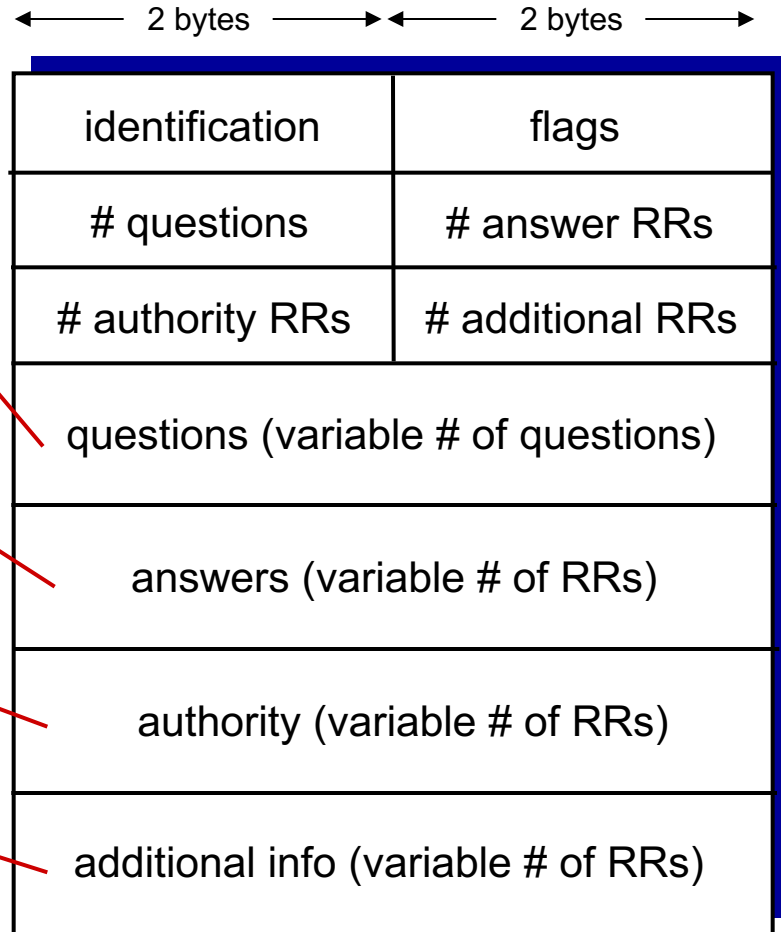
# DNS protocol, messages

Name & type fields  
(e.g., Type A or Type MX)

RRs in response  
to query  
(a reply can return  
multiple RRs)

records of other  
authoritative servers

additional “helpful”  
info that may be used



# DNS protocol, messages

For example, a reply to **an MX query**

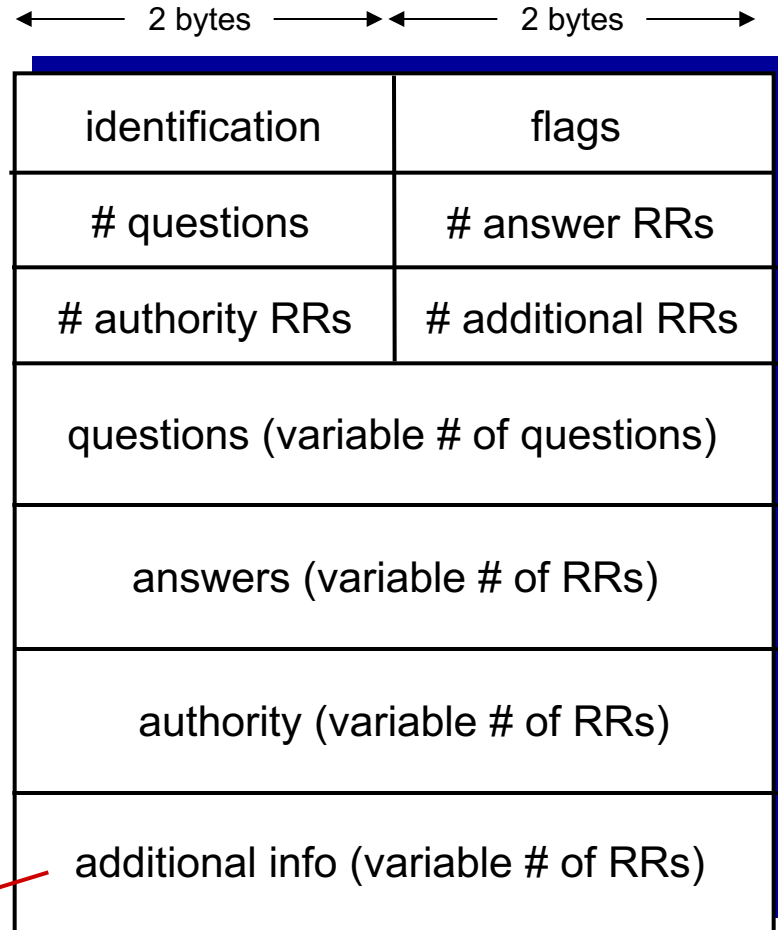
Answer section: Type MX

- an RR providing the canonical hostname of a mail server.

Additional section: Type A

- the IP address for the canonical hostname of the mail server.

additional “helpful”  
info that may be used



# DNS Overview

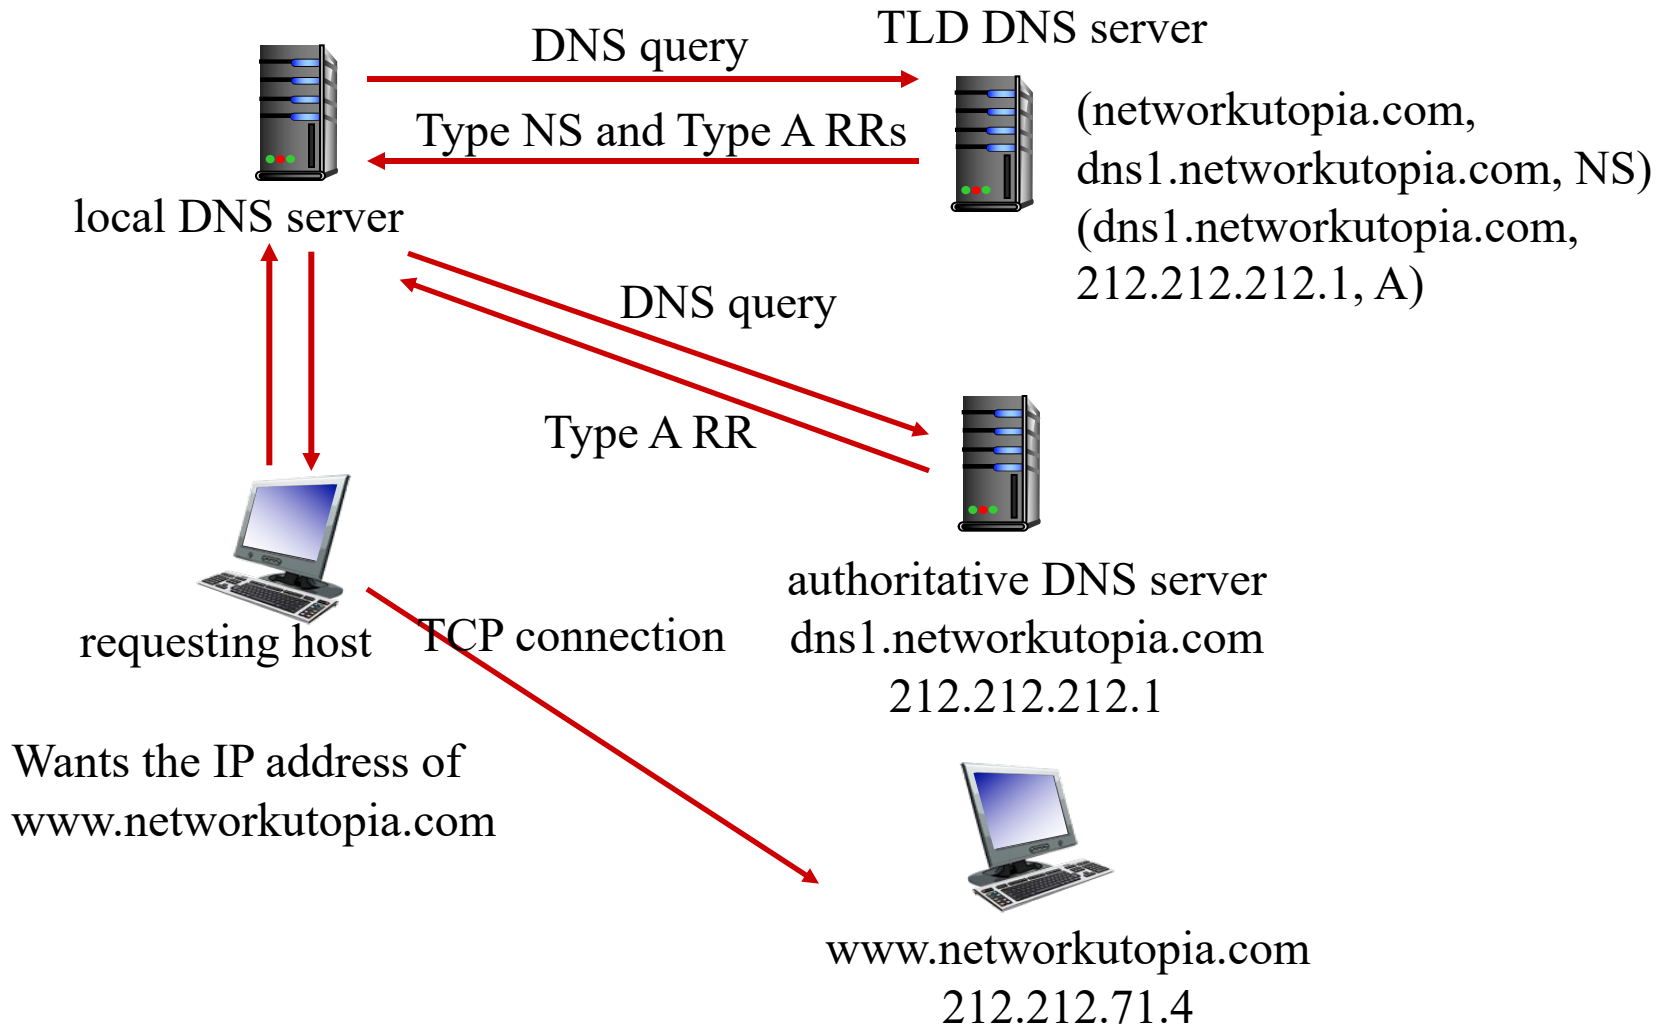
- DNS Services
- DNS Structure
  - Hierarchical structure
  - Iterated and recursive query
- DNS protocol
  - DNS Records
  - Query and reply messages
- Inserting records into DNS server

# Inserting records into DNS

- Example: new startup “Network Utopia”
- Register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
  - provide names, IP addresses of authoritative DNS server (primary and secondary)
  - registrar inserts two RRs into .com TLD server:  
(networkutopia.com, dns1.networkutopia.com, NS)  
(dns1.networkutopia.com, 212.212.212.1, A)



# Inserting records into DNS



# Attacking DNS

## Distributed denial-of-service (DDoS) attacks

- bombard root servers with traffic
  - not successful to date
  - traffic filtering
  - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
  - potentially more dangerous

## Redirect attacks

- man-in-middle
  - Intercept queries ; bogus reply
- DNS poisoning
  - Send bogus replies to DNS server

## Exploit DNS for DDoS

- target IP
- Redirect an unsuspecting Web user to attack Web site

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks


2.7 socket programming with UDP and TCP

# Pure P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

## Examples:

- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



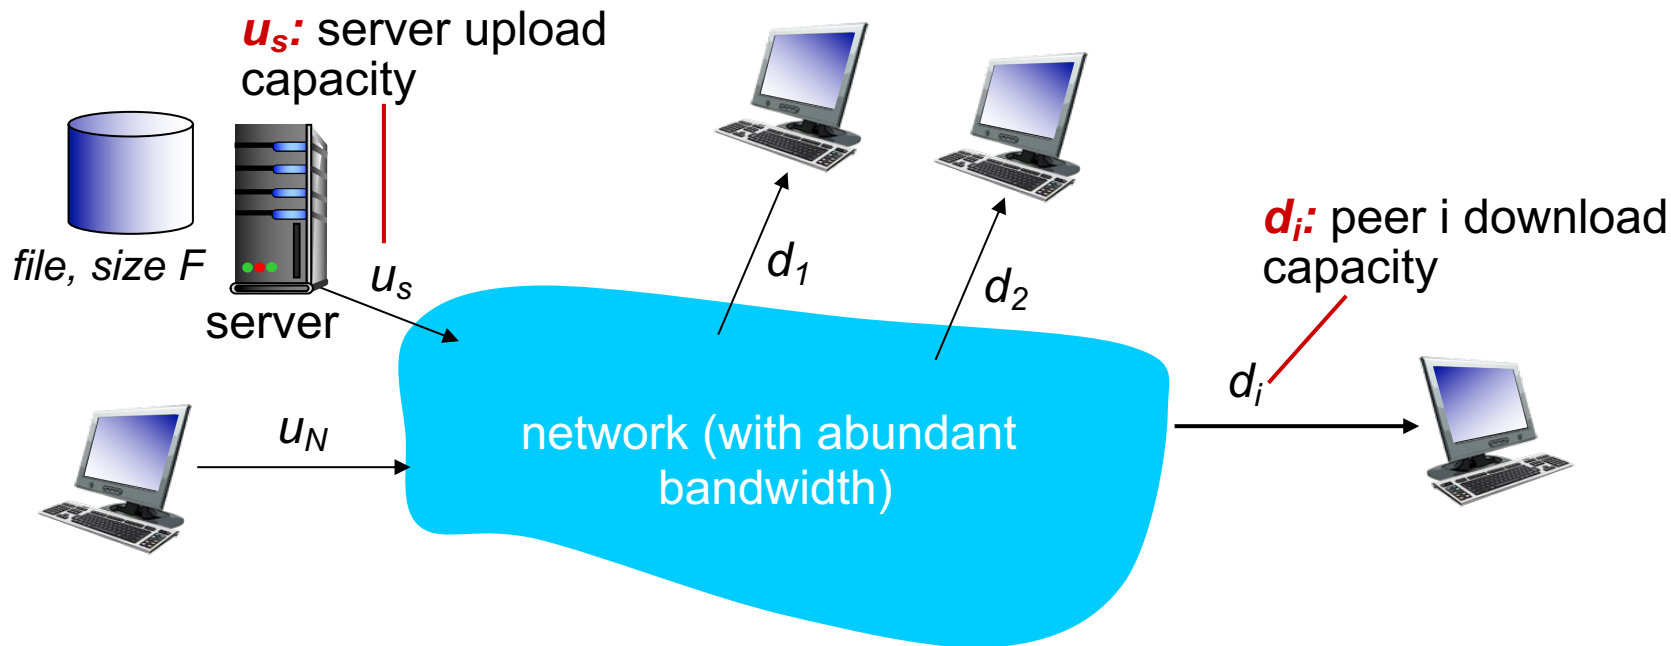
# DNS Overview

- P2P vs Client Server
- BitTorrent

# File distribution: client-server vs P2P

Question: How much time to distribute file (size  $F$ ) from one server to  $N$  peers?


- peer upload/download capacity is limited resource
- **Distribution time:** the time it takes to get a copy of the file to all  $N$  peers.



# File distribution time: client-server

- **Server transmission:** must sequentially send (upload)  $N$  file copies:

- time to send one copy:  $F/u_s$
- time to send  $N$  copies:  $NF/u_s$



- **Client:** each client must download file copy

- $d_{\min} = \min$  client download rate
- maximum client download time:  $F/d_{\min}$

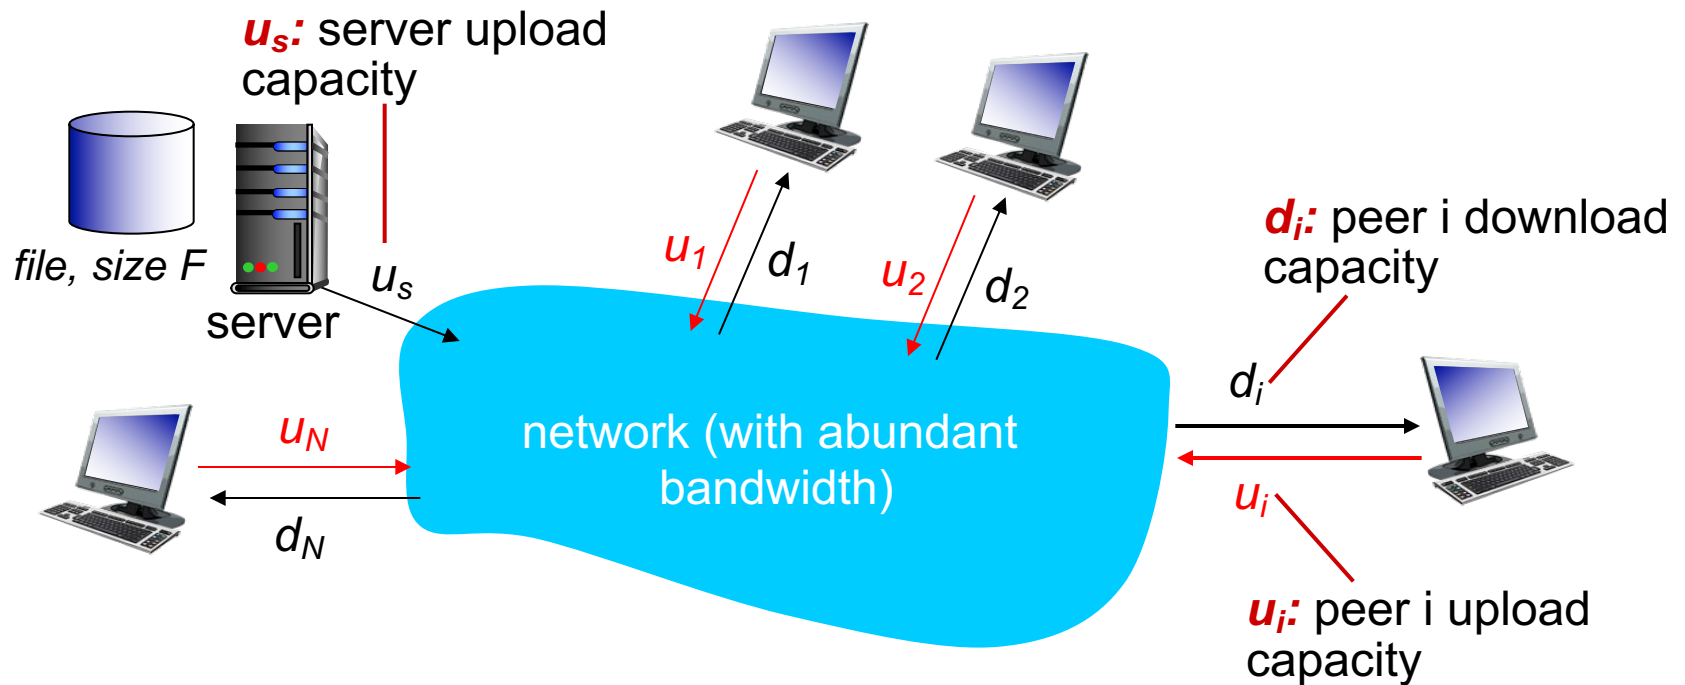
time to distribute  $F$   
to  $N$  clients using  
client-server approach

$$D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in  $N$

# File distribution time: P2P


In P2P model, clients are both downloaders and uploaders.





# File distribution time: P2P

- **Server transmission:** must upload at least one copy
  - time to send one copy:  $F/u_s$
- **Client downloading:** each client must download file copy
  - maximum client download time:  $F/d_{\min}$
- **Clients and server:** delivering a total of  $NF$  bits
  - max upload rate (limiting max download rate) is  $u_s + \sum u_i$



time to distribute  $F$   
to  $N$  clients using  
P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$


If each peer can redistribute a bit as soon as it receives the bit, then there is a scheme that actually achieves this lower bound

increases linearly in  $N$  ...

... but so does this, as each peer brings service capacity

# Client-server vs. P2P: example

client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$



# DNS Overview


- P2P vs Client Server
- BitTorrent

# P2P file distribution: BitTorrent

- File divided into 256Kb chunks
- Peers in BitTorrent send/receive file chunks


*tracker*: tracks peers participating in BitTorrent

*torrent*: group of peers exchanging chunks of a file



# P2P file distribution: BitTorrent

- Peer **joining** BitTorrent:
  - has no chunks, but will accumulate them over time from other peers
  - registers with tracker to get **list of peers**
  - TCP connections with subset of peers (“**neighbors**”)
- While **downloading**, peer uploads chunks to other peers
  - Peers may leave
  - Peers may come, initiating connections with Alice
- Once peer has entire file, it may (selfishly) leave or (altruistically) remain in BitTorrent



# BitTorrent: requesting, sending file chunks

Q1 : which chunks should she request first from her neighbors?

Q2: to which of her neighbors should she send requested chunks?

## requesting chunks:


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each “neighbor” for list of chunks that they have
- Alice requests missing chunks from peers, **rarest first**

## sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks **at highest rate**
  - other peers are choked by Alice (do not receive chunks from her)
  - re-evaluate every 10 secs
- every 30 secs: **randomly** select one **additional** peer, starts sending chunks
  - “optimistically unchoke” this peer
  - newly chosen peer may join top 4

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



*higher upload rate: find better trading partners, get file faster!*

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP